

## Problem Statement

Forecasting the number of views to the pages so that you can predict and optimize the ad placement for your clients.

We are provided with the data of 145k wikipedia pages and daily view count for each of them.

The clients belong to different regions and need data on how their ads will perform on pages in different languages.

```
In [1]: import pandas as pd
import numpy as np
import pylab as p
import matplotlib.pyplot as plot
import re
import os
import seaborn as sns
from collections import Counter
```

```
In [2]: import warnings
warnings.filterwarnings("ignore")
warnings.simplefilter("ignore")
```

```
In [3]: !gdown 1RPOLd_74kMrpkTg35X843hOrZIsSJSMF
```

Downloading...

From: [https://drive.google.com/uc?id=1RPOLd\\_74kMrpkTg35X843hOrZIsSJSMF](https://drive.google.com/uc?id=1RPOLd_74kMrpkTg35X843hOrZIsSJSMF) ([https://drive.google.com/uc?id=1RPOLd\\_74kMrpkTg35X843hOrZIsSJSMF](https://drive.google.com/uc?id=1RPOLd_74kMrpkTg35X843hOrZIsSJSMF))

To: /content/train\_1.csv

100% 278M/278M [00:04<00:00, 59.8MB/s]

```
In [4]: train = pd.read_csv('train_1.csv')
```

```
In [5]: train.head(10)
```

Out[5]:

|   | Page  | 2015-07-01 | 2015-07-02 | 2015-07-03 | 2015-07-04 | 2015-07-05 | 2015-07-06 | 2015-07-07 | 2015-07-08 | 2015-07-09 | ... | 2016-12-22 | 2016-12-23 | 2016-12-24 | 2016-12-25 |
|---|---|------------|------------|------------|------------|------------|------------|------------|------------|------------|-----|------------|------------|------------|------------|
| 0 | 2NE1_zh.wikipedia.org_all-access_spider           | 18.0       | 11.0       | 5.0        | 13.0       | 14.0       | 9.0        | 9.0        | 22.0       | 26.0       | ... | 32.0       | 63.0       | 15.0       | 26.0       |
| 1 | 2PM_zh.wikipedia.org_all-access_spider            | 11.0       | 14.0       | 15.0       | 18.0       | 11.0       | 13.0       | 22.0       | 11.0       | 10.0       | ... | 17.0       | 42.0       | 28.0       | 15.0       |
| 2 | 3C_zh.wikipedia.org_all-access_spider             | 1.0        | 0.0        | 1.0        | 1.0        | 0.0        | 4.0        | 0.0        | 3.0        | 4.0        | ... | 3.0        | 1.0        | 1.0        | 7.0        |
| 3 | 4minute_zh.wikipedia.org_all-access_spider        | 35.0       | 13.0       | 10.0       | 94.0       | 4.0        | 26.0       | 14.0       | 9.0        | 11.0       | ... | 32.0       | 10.0       | 26.0       | 27.0       |
| 4 | 52_Hz_I_Love_You_zh.wikipedia.org_all-access_s... | NaN        | NaN        | NaN        | NaN        | NaN        | NaN        | NaN        | NaN        | NaN        | ... | 48.0       | 9.0        | 25.0       | 13.0       |
| 5 | 5566_zh.wikipedia.org_all-access_spider           | 12.0       | 7.0        | 4.0        | 5.0        | 20.0       | 8.0        | 5.0        | 17.0       | 24.0       | ... | 16.0       | 27.0       | 8.0        | 17.0       |
| 6 | 91Days_zh.wikipedia.org_all-access_spider         | NaN        | NaN        | NaN        | NaN        | NaN        | NaN        | NaN        | NaN        | NaN        | ... | 2.0        | 7.0        | 33.0       | 8.0        |
| 7 | A'N'D_zh.wikipedia.org_all-access_spider          | 118.0      | 26.0       | 30.0       | 24.0       | 29.0       | 127.0      | 53.0       | 37.0       | 20.0       | ... | 64.0       | 35.0       | 35.0       | 28.0       |
| 8 | AKB48_zh.wikipedia.org_all-access_spider          | 5.0        | 23.0       | 14.0       | 12.0       | 9.0        | 9.0        | 35.0       | 15.0       | 14.0       | ... | 34.0       | 105.0      | 72.0       | 36.0       |
| 9 | ASCII_zh.wikipedia.org_all-access_spider          | 6.0        | 3.0        | 5.0        | 12.0       | 6.0        | 5.0        | 4.0        | 13.0       | 9.0        | ... | 25.0       | 17.0       | 22.0       | 29.0       |

10 rows × 551 columns

▬


▬



Checking for reasons for null values

```
In [6]: days = [r for r in range(1, len(train.columns))]  
plot.figure(figsize=(10,7))  
plot.xlabel('Day')  
plot.ylabel('Null values')  
plot.plot(days, train.isnull().sum()[1:])
```

Out[6]: [<matplotlib.lines.Line2D at 0x7d6e12fcc220>]



The null values are decreasing over time. So null values are caused because some pages were not created from the beginning.

### Treating null values

```
In [7]: train=train.dropna(how='all')
train=train.dropna(thresh=356)
```

```
In [8]: train=train.fillna(0)
train.tail()
```


Out[8]:

|        | Page  | 2015-07-01 | 2015-07-02 | 2015-07-03 | 2015-07-04 | 2015-07-05 | 2015-07-06 | 2015-07-07 | 2015-07-08 | 2015-07-09 | ... | 2016-12-22 | 2016-12-23 |
|--------|---|------------|------------|------------|------------|------------|------------|------------|------------|------------|-----|------------|------------|
| 144994 | Lucía_Hoyos_es.wikipedia.org_all-access_spider    | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | ... | 3.0        | 2.0        |
| 144995 | XXx_(película)_es.wikipedia.org_all-access_spider | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | ... | 6.0        | 4.0        |
| 144996 | ¿Volverías_con_tu_ex?_es.wikipedia.org_all-acc... | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | ... | 20.0       | 7.0        |
| 144997 | Alicia_a_través_del_espejo_(película)_es.wikip... | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | ... | 10.0       | 24.0       |
| 144999 | El_libro_de_la_selva_(película_de_2016)_es.wik... | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | 0.0        | ... | 11.0       | 11.0       |

5 rows × 551 columns

==

==



# EDA

```
In [9]: def split_page(page):  
        w = re.split('_|\\.', page)  
  
        return ' '.join(w[:-5]), w[-5], w[-2], w[-1]  
  
li = list(train.Page.apply(split_page))  
df = pd.DataFrame(li)  
df.columns = ['Title', 'Language', 'Access_type', 'Access_origin']  
df = pd.concat([train, df], axis = 1)
```

```
In [10]: df.head()
```

Out[10]:

|   | Page                                       | 2015-07-01 | 2015-07-02 | 2015-07-03 | 2015-07-04 | 2015-07-05 | 2015-07-06 | 2015-07-07 | 2015-07-08 | 2015-07-09 | ... | 2016-12-26 | 2016-12-27 | 2016-12-28 | 2016-12-29 | 2016-12-30 | 2016-12-31 |
|---|--|------------|------------|------------|------------|------------|------------|------------|------------|------------|-----|------------|------------|------------|------------|------------|------------|
| 0 | 2NE1_zh.wikipedia.org_all-access_spider    | 18.0       | 11.0       | 5.0        | 13.0       | 14.0       | 9.0        | 9.0        | 22.0       | 26.0       | ... | 14.0       | 20.0       | 22.0       | 19.0       | 18.0       | 20.0       |
| 1 | 2PM_zh.wikipedia.org_all-access_spider     | 11.0       | 14.0       | 15.0       | 18.0       | 11.0       | 13.0       | 22.0       | 11.0       | 10.0       | ... | 9.0        | 30.0       | 52.0       | 45.0       | 26.0       | 20.0       |
| 2 | 3C_zh.wikipedia.org_all-access_spider      | 1.0        | 0.0        | 1.0        | 1.0        | 0.0        | 4.0        | 0.0        | 3.0        | 4.0        | ... | 4.0        | 4.0        | 6.0        | 3.0        | 4.0        | 17.0       |
| 3 | 4minute_zh.wikipedia.org_all-access_spider | 35.0       | 13.0       | 10.0       | 94.0       | 4.0        | 26.0       | 14.0       | 9.0        | 11.0       | ... | 16.0       | 11.0       | 17.0       | 19.0       | 10.0       | 11.0       |
| 5 | 5566_zh.wikipedia.org_all-access_spider    | 12.0       | 7.0        | 4.0        | 5.0        | 20.0       | 8.0        | 5.0        | 17.0       | 24.0       | ... | 32.0       | 19.0       | 23.0       | 17.0       | 17.0       | 50.0       |

5 rows × 555 columns

==


==



```
In [13]: language_counts = df['Language'].value_counts()
```

```
# Create a count plot  
sns.countplot(data=df, x='Language')
```


```
Out[13]: <Axes: xlabel='Language', ylabel='count'>
```



```
In [15]: language_counts = df['Access_type'].value_counts()

# Create a count plot
sns.countplot(data=df, x='Access_type')
```

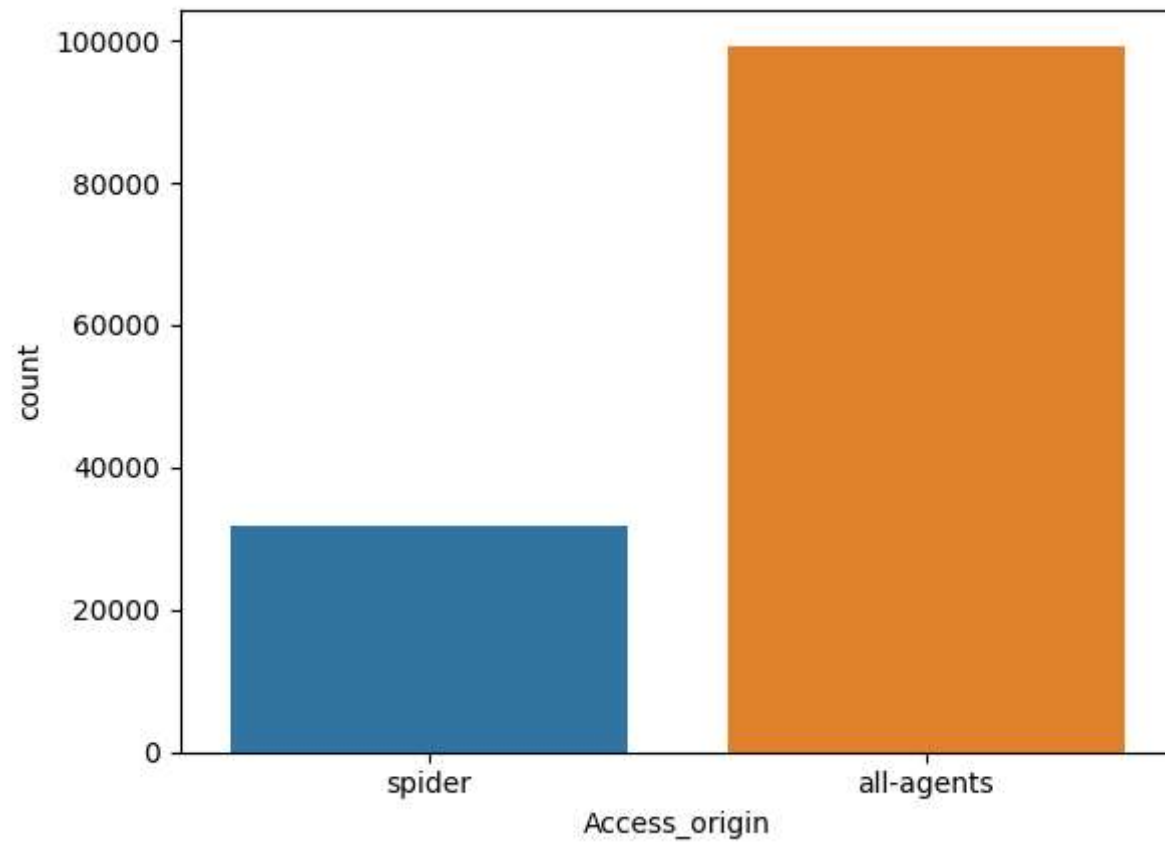
Out[15]: <Axes: xlabel='Access\_type', ylabel='count'>



```
In [17]: language_counts = df['Access_origin'].value_counts()

# Create a count plot
sns.countplot(data=df, x='Access_origin')
```

Out[17]: <Axes: xlabel='Access\_origin', ylabel='count'>





```
In [18]: df.groupby('Language').count()
```

Out[18]:

|          | Page  | 2015-07-01 | 2015-07-02 | 2015-07-03 | 2015-07-04 | 2015-07-05 | 2015-07-06 | 2015-07-07 | 2015-07-08 | 2015-07-09 | ... | 2016-12-25 | 2016-12-26 | 2016-12-27 | 2016-12-28 | 2016-12-29 | 2016-12-30 | 2016-12-31 |
|----------|-------|------------|------------|------------|------------|------------|------------|------------|------------|------------|-----|------------|------------|------------|------------|------------|------------|------------|
| Language |       |            |            |            |            |            |            |            |            |            |     |            |            |            |            |            |            |            |
| commons  | 7213  | 7213       | 7213       | 7213       | 7213       | 7213       | 7213       | 7213       | 7213       | 7213       | ... | 7213       | 7213       | 7213       | 7213       | 7213       | 7213       | 7213       |
| de       | 15505 | 15505      | 15505      | 15505      | 15505      | 15505      | 15505      | 15505      | 15505      | 15505      | ... | 15505      | 15505      | 15505      | 15505      | 15505      | 15505      | 15505      |
| en       | 19744 | 19744      | 19744      | 19744      | 19744      | 19744      | 19744      | 19744      | 19744      | 19744      | ... | 19744      | 19744      | 19744      | 19744      | 19744      | 19744      | 19744      |
| es       | 11516 | 11516      | 11516      | 11516      | 11516      | 11516      | 11516      | 11516      | 11516      | 11516      | ... | 11516      | 11516      | 11516      | 11516      | 11516      | 11516      | 11516      |
| fr       | 14630 | 14630      | 14630      | 14630      | 14630      | 14630      | 14630      | 14630      | 14630      | 14630      | ... | 14630      | 14630      | 14630      | 14630      | 14630      | 14630      | 14630      |
| ja       | 17172 | 17172      | 17172      | 17172      | 17172      | 17172      | 17172      | 17172      | 17172      | 17172      | ... | 17172      | 17172      | 17172      | 17172      | 17172      | 17172      | 17172      |
| ru       | 12329 | 12329      | 12329      | 12329      | 12329      | 12329      | 12329      | 12329      | 12329      | 12329      | ... | 12329      | 12329      | 12329      | 12329      | 12329      | 12329      | 12329      |
| www      | 5530  | 5530       | 5530       | 5530       | 5530       | 5530       | 5530       | 5530       | 5530       | 5530       | ... | 5530       | 5530       | 5530       | 5530       | 5530       | 5530       | 5530       |
| zh       | 14187 | 14187      | 14187      | 14187      | 14187      | 14187      | 14187      | 14187      | 14187      | 14187      | ... | 14187      | 14187      | 14187      | 14187      | 14187      | 14187      | 14187      |

9 rows × 554 columns



```
In [19]: #here we see that the Languages are not treated properly as there are commons and www
```

In [20]: *# Checking another way of fetching the Language out of the string*

```
def lang(Page):
    val = re.search('[a-z][a-z].wikipedia.org',Page)
    if val:
        return val[0][0:2]

    return 'no_lang'

df['Language']=df['Page'].apply(lambda x: lang(str(x)))
```

In [21]: *df.groupby('Language').count() #now the count has increased. You can go back and get it sorted*

Out[21]:

|          | Page  | 2015-07-01 | 2015-07-02 | 2015-07-03 | 2015-07-04 | 2015-07-05 | 2015-07-06 | 2015-07-07 | 2015-07-08 | 2015-07-09 | ... | 2016-12-25 | 2016-12-26 | 2016-12-27 | 2016-12-28 | 2016-12-29 | 2016-12-30 | 2016-12-31 |
|----------|-------|------------|------------|------------|------------|------------|------------|------------|------------|------------|-----|------------|------------|------------|------------|------------|------------|------------|
| Language |       |            |            |            |            |            |            |            |            |            |     |            |            |            |            |            |            |            |
| de       | 17123 | 17123      | 17123      | 17123      | 17123      | 17123      | 17123      | 17123      | 17123      | 17123      | ... | 17123      | 17123      | 17123      | 17123      | 17123      | 17123      | 17         |
| en       | 21938 | 21938      | 21938      | 21938      | 21938      | 21938      | 21938      | 21938      | 21938      | 21938      | ... | 21938      | 21938      | 21938      | 21938      | 21938      | 21938      | 21         |
| es       | 13397 | 13397      | 13397      | 13397      | 13397      | 13397      | 13397      | 13397      | 13397      | 13397      | ... | 13397      | 13397      | 13397      | 13397      | 13397      | 13397      | 13         |
| fr       | 16717 | 16717      | 16717      | 16717      | 16717      | 16717      | 16717      | 16717      | 16717      | 16717      | ... | 16717      | 16717      | 16717      | 16717      | 16717      | 16717      | 16         |
| ja       | 19115 | 19115      | 19115      | 19115      | 19115      | 19115      | 19115      | 19115      | 19115      | 19115      | ... | 19115      | 19115      | 19115      | 19115      | 19115      | 19115      | 19         |
| no_lang  | 13856 | 13856      | 13856      | 13856      | 13856      | 13856      | 13856      | 13856      | 13856      | 13856      | ... | 13856      | 13856      | 13856      | 13856      | 13856      | 13856      | 13         |
| ru       | 14109 | 14109      | 14109      | 14109      | 14109      | 14109      | 14109      | 14109      | 14109      | 14109      | ... | 14109      | 14109      | 14109      | 14109      | 14109      | 14109      | 14         |
| zh       | 14791 | 14791      | 14791      | 14791      | 14791      | 14791      | 14791      | 14791      | 14791      | 14791      | ... | 14791      | 14791      | 14791      | 14791      | 14791      | 14791      | 14         |

8 rows × 554 columns

==

==



```
In [22]: df_language=df.groupby('Language').mean().transpose()
df_language
```

```
Out[22]:
```

| Language   | de          | en          | es          | fr         | ja          | no_lang    | ru          | zh         |
|------------|-------------|-------------|-------------|------------|-------------|------------|-------------|------------|
| 2015-07-01 | 774.426269  | 3861.418270 | 1140.148541 | 505.989352 | 620.425007  | 107.428551 | 670.765965  | 280.230884 |
| 2015-07-02 | 763.878176  | 3848.940651 | 1089.566843 | 509.238500 | 712.459639  | 112.593533 | 682.374796  | 280.646745 |
| 2015-07-03 | 733.166910  | 3654.264017 | 1002.042323 | 489.659688 | 643.454355  | 106.425664 | 632.463605  | 278.785275 |
| 2015-07-04 | 672.798867  | 3804.480354 | 940.835112  | 523.400191 | 808.439236  | 90.827223  | 594.882132  | 281.481036 |
| 2015-07-05 | 782.125153  | 3929.175312 | 1023.180190 | 513.872346 | 775.587654  | 100.659353 | 633.532072  | 300.263944 |
| ...        | ...         | ...         | ...         | ...        | ...         | ...        | ...         | ...        |
| 2016-12-27 | 1123.751153 | 6415.150971 | 1078.762559 | 846.143028 | 808.217578  | 161.744515 | 1007.479977 | 364.880738 |
| 2016-12-28 | 1066.064241 | 6208.003464 | 1117.616705 | 774.151223 | 807.662046  | 186.088770 | 953.722801  | 371.549185 |
| 2016-12-29 | 1036.503358 | 6630.178002 | 1067.152571 | 757.718610 | 885.257233  | 157.015156 | 917.736339  | 342.295788 |
| 2016-12-30 | 986.052678  | 5490.552740 | 813.258789  | 712.405097 | 981.830133  | 162.524682 | 822.712028  | 344.481847 |
| 2016-12-31 | 938.918939  | 5367.033777 | 782.740240  | 658.079739 | 1230.408789 | 140.588265 | 911.081579  | 352.220337 |

550 rows × 8 columns


==

==

```
In [23]: df_language.reset_index(inplace=True)
df_language.set_index('index', inplace=True)
```

```
In [24]: df_language.plot(figsize=(12,7))  
         plot.ylabel('Views per Page')
```


Out[24]: Text(0, 0.5, 'Views per Page')



```
In [25]: ## English gets most amount of views followed by spanish
```

```
In [26]: # printing English
df_language['en'].plot(figsize=(12,7))
plot.ylabel('Views per Page')
```

Out[26]: Text(0, 0.5, 'Views per Page')



```
In [27]: total_view=df_language.copy()
```

## Checking the stationarity


```
In [28]: # Dickey-Fuller test
```

```
In [29]: from statsmodels.tsa.stattools import adfuller
def df_test(x):
    result=adfuller(x)
    print('ADF Stastistic: %f'%result[0])
    print('p-value: %f'%result[1])

df_test(total_view['en'])
```

ADF Stastistic: -2.403786  
p-value: 0.140671

```
In [30]: # We see that the p value is not low enough(<0.05). Therefore, we can say our series in not stationary as we
```



## Making the time series stationary

```
In [31]: ts=total_view[['en']]
```

```
In [32]: ts = ts.reset_index()
```

```
In [33]: ts.index = pd.to_datetime(ts.index)
```

```
In [34]: ts.set_index('index', inplace=True)
```

In [35]:

```
ts
```

Out[35]:


| Language   | en          |
|------------|-------------|
| index      |             |
| 2015-07-01 | 3861.418270 |
| 2015-07-02 | 3848.940651 |
| 2015-07-03 | 3654.264017 |
| 2015-07-04 | 3804.480354 |
| 2015-07-05 | 3929.175312 |
| ...        | ...         |
| 2016-12-27 | 6415.150971 |
| 2016-12-28 | 6208.003464 |
| 2016-12-29 | 6630.178002 |
| 2016-12-30 | 5490.552740 |
| 2016-12-31 | 5367.033777 |

550 rows × 1 columns

==

==


```
In [36]: import statsmodels.api as sm
model = sm.tsa.seasonal_decompose(ts, period=7) # Knowing the right period is V.Imp and You can use ACF to g
model.plot();
```





```
In [37]: model.resid.plot()
```

```
Out[37]: <Axes: xlabel='index'>
```




```
In [38]: df_test(model.resid.dropna())
```

```
ADF Stastistic: -11.435041  
p-value: 0.000000
```

```
In [39]: ## We can see that our series is now stationary, we can also try differencing to see what results we can get.
```

## Remove trend and seasonality with differencing

```
In [40]: ts_diff = ts - ts.shift(1)
         plot.plot(ts_diff.values)
         plot.show()
```




```
In [41]: ts_diff.dropna(inplace=True)
         df_test(ts_diff)
```


ADF Stastistic: -8.261352  
p-value: 0.000000

In [42]: *#Also the  $p$  value is 0. So we can say that our graph is now stationery. Now we can apply the ARIMA model*

## Plot the autocorreltaion and partial auto correlation functions

```
In [43]: from statsmodels.graphics.tsaplots import plot_acf, plot_pacf  
acf=plot_acf(ts_diff,lags=20)  
pacf=plot_pacf(ts_diff,lags=20)
```





```
In [44]: ts.diff(7) #
```

```
Out[44]:
```

| Language   | en          |
|------------|-------------|
| index      |             |
| 2015-07-01 | NaN         |
| 2015-07-02 | NaN         |
| 2015-07-03 | NaN         |
| 2015-07-04 | NaN         |
| 2015-07-05 | NaN         |
| ...        | ...         |
| 2016-12-27 | 804.087291  |
| 2016-12-28 | 997.424970  |
| 2016-12-29 | 1346.841371 |
| 2016-12-30 | 561.319309  |
| 2016-12-31 | 120.887136  |

550 rows × 1 columns

==

==

## ARIMA MODEL

```
In [45]: from statsmodels.tsa.statespace.sarimax import SARIMAX
from pandas import DataFrame
```

```
In [46]: model = SARIMAX(ts, order=(1,1,1))
        model_fit = model.fit(dis=0)
```


```
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
    self._init_dates(dates, freq)
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
    self._init_dates(dates, freq)
```

```
In [47]: model_fit.predict(12)
```

```
Out[47]: 2015-07-13    4210.297557
        2015-07-14    4174.226495
        2015-07-15    4197.775650
        2015-07-16    3902.374198
        2015-07-17    3945.629070
        ...
        2016-12-27    6062.278727
        2016-12-28    6317.880194
        2016-12-29    6130.967200
        2016-12-30    6516.814683
        2016-12-31    5480.984400
        Freq: D, Name: predicted_mean, Length: 538, dtype: float64
```

```
In [48]: train = ts[:-20]
        test = ts[-20:]
```

```
In [49]: ## In ARMA family models, for getting p (Auto-regressive order) we look at PACF plot, there is a cut-off at l
```



```
In [50]: model = SARIMAX(train, order=(2, 1, 2))
         fitted = model.fit(dispatch=-1)

         # Forecast
         fc = fitted.forecast(20, alpha=0.02)

         # Make as pandas series
         fc_series = pd.Series(fc, index=test.index)
         # Plot
         plot.figure(figsize=(12,5), dpi=100)
         plot.plot(train, label='training')
         plot.plot(test, label='actual')
         plot.plot(fc_series, label='forecast')

         plot.title('Forecast vs Actuals')
         plot.legend(loc='upper left', fontsize=8)
```

```
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
```

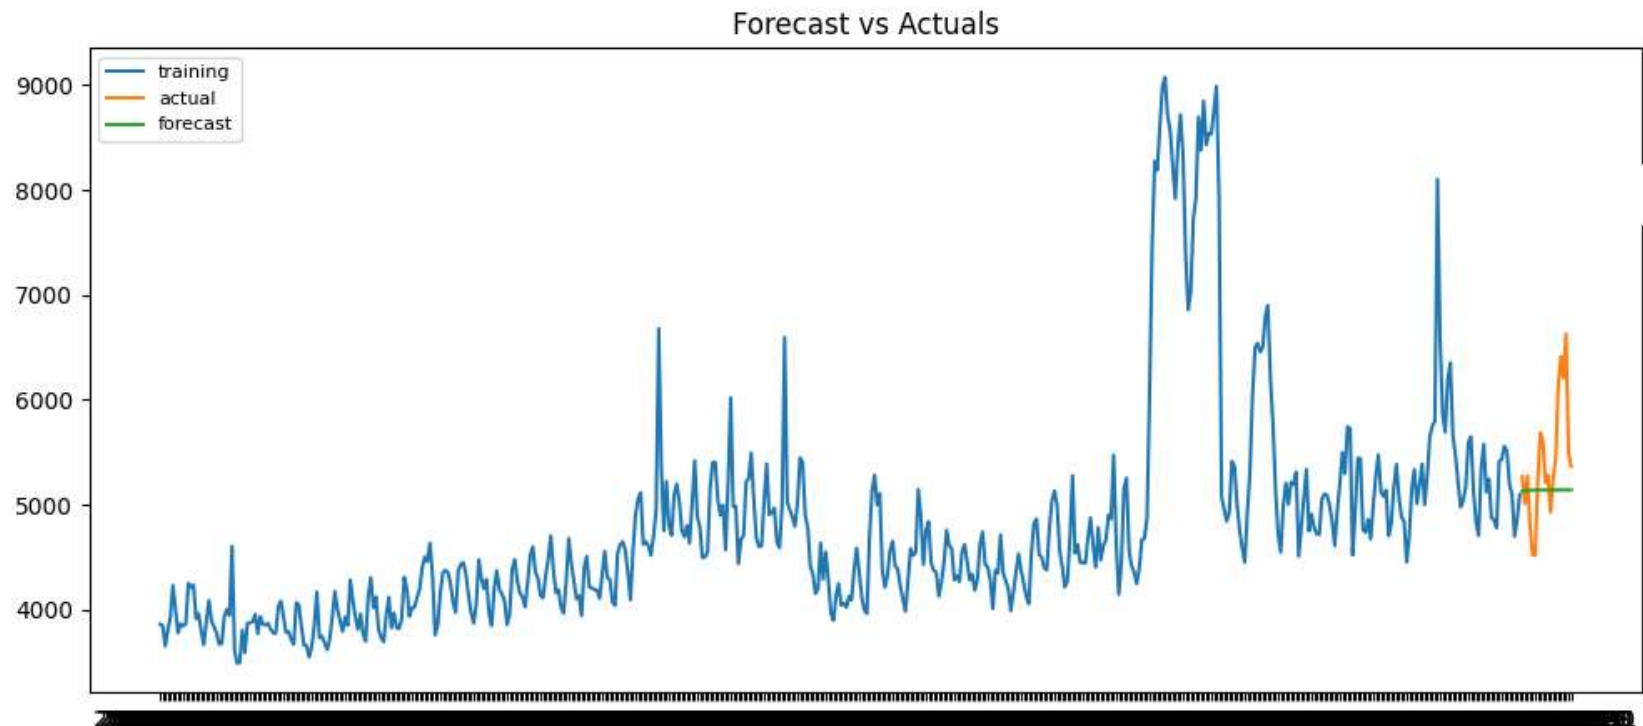
```
    self._init_dates(dates, freq)
```

```
/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.
```

```
    self._init_dates(dates, freq)
```

```
Out[50]: <matplotlib.legend.Legend at 0x7d6db1697130>
```





```
In [51]: mape = np.mean(np.abs(fc.values - test.values)/np.abs(test.values))  
rmse = np.mean((fc.values - test.values)**2)**.5  
print("mape:", mape)  
print("rsme:", rmse)
```

```
mape: 0.08213219335167705  
rsme: 624.4451158083979
```

## using exogenous data

```
In [52]: ex_df = pd.read_csv('Exog_Campaign_eng.csv')  
ex_df.head()
```

```
Out[52]:
```

|   | Exog |
|---|------|
| 0 | 0    |
| 1 | 0    |
| 2 | 0    |
| 3 | 0    |
| 4 | 0    |

==  
==

```
In [53]: exog=ex_df['Exog'].to_numpy()
```

```

In [54]: import statsmodels.api as sm
train=ts[:530]
test=ts[530:]
model=sm.tsa.statespace.SARIMAX(train,order=(1, 1, 1),seasonal_order=(1,1,1,7),exog=exog[:530])
results=model.fit()

fc=results.forecast(20,dynamic=True,exog=pd.DataFrame(exog[530:]))

# Make as pandas series
fc_series = pd.Series(fc)
# Plot
train.index=train.index.astype('datetime64[ns]')
test.index=test.index.astype('datetime64[ns]')
plot.figure(figsize=(12,5), dpi=100)
plot.plot(train, label='training')
plot.plot(test, label='actual')
plot.plot(fc_series, label='forecast')

plot.title('Forecast vs Actuals')
plot.legend(loc='upper left', fontsize=8)

```

```

/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.

```

```

    self._init_dates(dates, freq)

```

```

/usr/local/lib/python3.10/dist-packages/statsmodels/tsa/base/tsa_model.py:473: ValueWarning: No frequency information was provided, so inferred frequency D will be used.

```

```


    self._init_dates(dates, freq)

```

```

Out[54]: <matplotlib.legend.Legend at 0x7d6daed52080>

```



```
In [55]: mape = np.mean(np.abs(fc.values - test.values)/np.abs(test.values))  
rmse = np.mean((fc.values - test.values)**2)**.5  
print("mape:", mape)  
print("rsme:", rmse)
```

```
mape: 0.10347693827344548  
rsme: 737.5231800914678
```

## Facebook Prophet

```
In [56]: # !pip install pystan~=2.14  
# !pip install fbprophet
```

```
In [57]: ts_df = ts.reset_index().copy()
ts_df.columns = [['ds', 'y']]
```

```
In [58]: df2=ts_df.copy()
df2['exog'] = exog
df2.columns = ['ds', 'y', 'exog']
df2.head()
```

```
Out[58]:
```

|   |            | ds          | y | exog |
|---|------------|-------------|---|------|
| 0 | 2015-07-01 | 3861.418270 |   | 0    |
| 1 | 2015-07-02 | 3848.940651 |   | 0    |
| 2 | 2015-07-03 | 3654.264017 |   | 0    |
| 3 | 2015-07-04 | 3804.480354 |   | 0    |
| 4 | 2015-07-05 | 3929.175312 |   | 0    |

=====

=====

```
In [59]: df2[:20].info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 530 entries, 0 to 529
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype  
---  -
0    ds      530 non-null     object 
1    y        530 non-null     float64
2    exog     530 non-null     int64  
dtypes: float64(1), int64(1), object(1)
memory usage: 12.5+ KB
```

```
In [60]: from prophet import Prophet
m = Prophet(weekly_seasonality=True)
m.fit(df2[['ds', 'y']][:20])
future = m.make_future_dataframe(periods=20, freq="D")
forecast = m.predict(future)
fig = m.plot(forecast)
```

INFO:prophet:Disabling yearly seasonality. Run prophet with yearly\_seasonality=True to override this.

INFO:prophet:Disabling daily seasonality. Run prophet with daily\_seasonality=True to override this.

DEBUG:cmdstanpy:input tempfile: /tmp/tmpbugd82zr/bihxspqn.json

DEBUG:cmdstanpy:input tempfile: /tmp/tmpbugd82zr/ewtm9c\_1.json

DEBUG:cmdstanpy:idx 0

DEBUG:cmdstanpy:running CmdStan, num\_threads: None


DEBUG:cmdstanpy:CmdStan args: ['/usr/local/lib/python3.10/dist-packages/prophet/stan\_model/prophet\_model.bin', 'random', 'seed=19574', 'data', 'file=/tmp/tmpbugd82zr/bihxspqn.json', 'init=/tmp/tmpbugd82zr/ewtm9c\_1.json', 'output', 'file=/tmp/tmpbugd82zr/prophet\_model192hqy3a1/prophet\_model-20230928175519.csv', 'method=optimize', 'algorithm=lbfgs', 'iter=10000']

17:55:19 - cmdstanpy - INFO - Chain [1] start processing

INFO:cmdstanpy:Chain [1] start processing

17:55:19 - cmdstanpy - INFO - Chain [1] done processing

INFO:cmdstanpy:Chain [1] done processing



```
In [61]: model2=Prophet(interval_width=0.9, weekly_seasonality=True, changepoint_prior_scale=1)
model2.add_regressor('exog')
model2.fit(df2[:-20])
forecast2 = model2.predict(df2)
fig = model2.plot(forecast2)
```

INFO:prophet:Disabling yearly seasonality. Run prophet with yearly\_seasonality=True to override this.

INFO:prophet:Disabling daily seasonality. Run prophet with daily\_seasonality=True to override this.

DEBUG:cmdstanpy:input tempfile: /tmp/tmpbugd82zr/b44qqhs1.json

DEBUG:cmdstanpy:input tempfile: /tmp/tmpbugd82zr/6mz6nkr4.json

DEBUG:cmdstanpy:idx 0

DEBUG:cmdstanpy:running CmdStan, num\_threads: None

DEBUG:cmdstanpy:CmdStan args: ['/usr/local/lib/python3.10/dist-packages/prophet/stan\_model/prophet\_model.bin', 'random', 'seed=63734', 'data', 'file=/tmp/tmpbugd82zr/b44qqhs1.json', 'init=/tmp/tmpbugd82zr/6mz6nkr4.json', 'output', 'file=/tmp/tmpbugd82zr/prophet\_model\_ejkvmlw/prophet\_model-20230928175519.csv', 'method=optimize', 'algorithm=lbfgs', 'iter=10000']


17:55:19 - cmdstanpy - INFO - Chain [1] start processing

INFO:cmdstanpy:Chain [1] start processing

17:55:20 - cmdstanpy - INFO - Chain [1] done processing


INFO:cmdstanpy:Chain [1] done processing





```
In [62]: y_true = df2['y'].values
y_pred = forecast2['yhat'].values

plot.plot(y_true, label='Actual')
plot.plot(y_pred, label='Predicted')
plot.legend()
plot.show()
```



```
In [63]: mape = np.mean(np.abs(forecast2['yhat'][-20:] - df2['y'][-20:].values)/np.abs(df2['y'][-20:].values))
print("mape:", mape)
```

mape: 0.06757807513038254

# Questionnaire

## 1. Defining the problem statements and where can this and modifications of this be used?

Forecasting the number of views to the pages so that you can predict and optimize the ad placement for your clients.

The same can be used for predicting the future sale in a online store, predicting the using pattern of streaming app etc.

## 2. Write 3 inferences you made from the data visualizations

The amount of null values decreasing over the time. This is because some website were not there from start

There the more number of pages in English Language

The desktop and mobile user are almost similar

## 3. What does the decomposition of series do?

Decomposition remove trend and seasonlity from the time series data.

## 4. What level of differencing gave you a stationary series?

1 level of differencing gave a stationary data

## 5. Difference between arima, sarima & sarimax.

ARIMA is a basic time series forecasting model that consist of three main component; AutoRegressive (AR), Moving Average (MA) and Residual

SARIMA is an extension of ARIMA model that takes into account easonlity in time series data

SARIMAX is an extensioin of SARIMA that allows for the inclusion of exogenous variable which are external factors that can influence the time series

## 6. Compare the number of views in different languages

Large number of viewers are English speaking

Followed by Chinese and Spanish

## 7. What other methods other than grid search would be suitable to get the model for all languages?

Random search

Bayesian Optimization

In [63]: